

OK, thus far we have been talking about linear models. All these can be viewed as a single-layer neural net. The next step is to move on to multi-layer nets. Training these is a bit more involved, and implementing from scratch requires time and effort. Instead, we just use well-established libraries. I prefer PyTorch, which is based on an earlier library called Torch (designed for training neural nets via backprop).

NAME : Y.HEMANTH TEJA EMAIL : hy1713@nyu.edu

```
import numpy as np
import torch
import torchvision

import matplotlib.pyplot as plt
import random
import pandas as pd
import sklearn
```

Torch handles data types a bit differently. Everything in torch is a *tensor*.

```
a = np.random.rand(2,3)
b = torch.from_numpy(a)

# Q4.1 Display the contents of a, b
# ...
# ...
print("Displaying contents of A")
print(a)

print("\nDisplaying contents of B")
print(b)

[> Displaying contents of A
[[0.86780681 0.45833768 0.55640672]
 [0.83049916 0.74459275 0.91870588]]

Displaying contents of B
tensor([[0.8678, 0.4583, 0.5564],
        [0.8305, 0.7446, 0.9187]], dtype=torch.float64)
```

The idea in Torch is that tensors allow for easy forward (function evaluations) and backward (gradient) passes.

```
A = torch.rand(2,2)
b = torch.rand(2,1)
x = torch.rand(2,1, requires_grad=True)
```

```

x = torch.rand(2,1, requires_grad=True)

y = torch.matmul(A,x) + b

print(y)
z = y.sum()
print(z)
z.backward()
print(x.grad)
print(x)

[>] tensor([[1.3683],
           [1.3693]], grad_fn=<AddBackward0>)
      tensor(2.7376, grad_fn=<SumBackward0>)
      tensor([[1.4168],
           [1.1964]])
      tensor([[0.8909],
           [0.5900]], requires_grad=True)

```

Notice how the backward pass computed the gradients using autograd. OK, enough background. Time to train some networks. Let us load the *Fashion MNIST* dataset, which is a database of grayscale images of clothing items.

```

trainingdata = torchvision.datasets.FashionMNIST('./FashionMNIST/',train=True,download=True)
testdata = torchvision.datasets.FashionMNIST('./FashionMNIST/',train=False,download=True)

```

Let us examine the size of the dataset.

```

# Q4.2 How many training and testing data points are there in the dataset?
# What is the number of features in each data point?
#
# ...

```

```

print(trainingdata)
print("\nTotal number of training data points = 60000\n")
print(testdata)
print("\nTotal number of testing data points is 10000")

```

```
[>]
```

```
Dataset FashionMNIST
  Number of datapoints: 60000
  Root location: ./FashionMNIST/
  Split: Train
  StandardTransform
  Transform: ToTensor()
```

```
Total number of training data points = 60000
```

```
len(trainingdata)
```

```
↳ 60000
```

```
Split: Test
```

```
len(testdata)
```

```
↳ 10000
```


Let us try to visualize some of the images. Since each data point is a tensor (not an array) we need to postprocess to use matplotlib.

```
import matplotlib.pyplot as plt
%matplotlib inline
```

```
image, label = trainingdata[0]
# Q4.3 Assuming each sample is an image of size 28x28, show it in matplotlib.
# ...
# ...
```

```
plt.imshow(image[0].numpy().squeeze(), cmap = 'Greys_r')
```

```
↳ <matplotlib.image.AxesImage at 0x7fee10464278>
```



Let's try plotting several images. This is conveniently achieved in PyTorch using a *data loader*, which loads data in batches.

```
trainDataLoader = torch.utils.data.DataLoader(trainingdata, batch size=64, shuffle=True)
```

```

testDataLoader = torch.utils.data.DataLoader(testdata, batch_size=64, shuffle=False)
images, labels = iter(trainDataLoader).next()
print(images.size(), labels)

[> torch.Size([64, 1, 28, 28]) tensor([1, 3, 7, 4, 9, 8, 0, 8, 8, 3, 6, 1, 9, 9, 7,
      0, 5, 4, 4, 8, 0, 6, 6, 0, 4, 0, 4, 7, 4, 0, 7, 9, 2, 9, 4, 9, 3, 5, 7,
      7, 2, 7, 3, 0, 9, 5, 5, 1, 7, 8, 1, 7, 1, 5, 7])]

def output_label(label):
    output_mapping = {
        0: "T-shirt/Top",
        1: "Trouser",
        2: "Pullover",
        3: "Dress",
        4: "Coat",
        5: "Sandal",
        6: "Shirt",
        7: "Sneaker",
        8: "Bag",
        9: "Ankle Boot"
    }
    input = (label.item() if type(label) == torch.Tensor else label)
    return output_mapping[input]

# Q4.4 Visualize the first 10 images of the first minibatch
# returned by testDataLoader.
# ...
# ...

printOutput = torch.utils.data.DataLoader(testdata, batch_size=10)
batch = next(iter(printOutput))
images, labels = batch
print(type(images), type(labels))
print(images.shape, labels.shape)

[> <class 'torch.Tensor'> <class 'torch.Tensor'>
      torch.Size([10, 1, 28, 28]) torch.Size([10])

grid = torchvision.utils.make_grid(images, nrow=10)
plt.figure(figsize=(15, 20))
plt.imshow(np.transpose(grid, (1, 2, 0)))
print("labels: ", end=" ")
for i, label in enumerate(labels):
    print(output_label(label), end=" ")

[>

```

labels: Ankle Boot, Pullover, Trouser, Trouser, Shirt, Trouser, Coat, Shirt, Sai



Now we are ready to define our linear model. Here is some boilerplate PyTorch code that implements the forward model for a single layer network for logistic regression (similar to the one discussed in class notes).

```
class LinearReg(torch.nn.Module):
    def __init__(self, input_dim, output_dim):
        super(LinearReg, self).__init__()
        self.linear = torch.nn.Linear(input_dim, output_dim)

    def forward(self, x):
        x = x.view(-1, 28*28)
        transformed_x = self.linear(x)
        return transformed_x
```

```
input_dim = 28*28
output_dim = 10
```

```
model = LinearReg(input_dim, output_dim)
```

```
Loss = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.03)
```

```
len(trainDataLoader)
```

```
↳ 938
```

```
print(model)
print(model.parameters())
print(len(list(model.parameters())))
```

```
↳ LinearReg(
  (linear): Linear(in_features=784, out_features=10, bias=True)
)
<generator object Module.parameters at 0x7fee10034e08>
2
```

Cool! Now we have set everything up. Let's try to train the network.

```
# Q4.5 Write down a for-loop that trains this network for 20 minibatch iterations,
# and print the train/test losses.
# Save them in the variables above. If done correctly, you should be able to
# execute the next code block.
# ...
```

```

# ...
train_loss_history = []
test_loss_history = []
for i, (images, labels) in enumerate(trainDataLoader):

    images = images.view(-1, 28*28).requires_grad_()
    labels = labels

    # Clear gradients with respect to the parameters
    optimizer.zero_grad()

    # Forward pass
    outputs = model(images)

    # Calculate Loss
    calcLoss = Loss(outputs, labels)

    # Getting gradients with respect to the parameters
    calcLoss.backward()

    #Update the parameters
    optimizer.step()
    train_loss_history.append(calcLoss.item())
    print(calcLoss.item())

    correctPredictions = 0
    totalNumLabels = 0

    if(i >=19):
        break

for i, (images, labels) in enumerate(testDataLoader):

    images = images.view(-1, 28*28).requires_grad_()

    outputs = model(images)
    calcLoss = Loss(outputs, labels)

    if(i%500 == 0):
        test_loss_history.append(calcLoss.item())

    # Get predictions from the maximum value
    _, predicted = torch.max(outputs.data, 1)

    #Figuring the total number of labels
    totalNumLabels += labels.size(0)

    # Figuring the total number of correct predictions
    correctPredictions += (predicted == labels).sum()

    #Computing accuracy
    accuracy = 100 * correctPredictions.numpy() / totalNumLabels


```

```
accuracy = 100 * correct_predictions.numpy() / total_train_labels
```

```
↳ 2.325894355773926  
2.247469663619995  
2.20113205909729  
2.0652406215667725  
2.0560007095336914  
2.0687193870544434  
1.9734618663787842  
1.8722034692764282  
1.8808752298355103  
1.9147193431854248  
1.7352871894836426  
1.8906794786453247  
1.7159597873687744  
1.7664744853973389  
1.74052095413208  
1.6854078769683838  
1.5379551649093628  
1.6379114389419556  
1.6449737548828125  
1.5461499691009521
```

```
plt.plot(train_loss_history, '-', linewidth=3, label='Train error')  
plt.xlabel('epoch')  
plt.ylabel('loss')  
plt.grid(True)  
plt.legend()
```


```
↳ <matplotlib.legend.Legend at 0x7fee101bc0b8>
```



```
plt.plot(test_loss_history, '-', linewidth=3, label='Test error')  
plt.xlabel('epoch')  
plt.ylabel('loss')  
plt.grid(True)  
plt.legend()
```

```
↳
```

<matplotlib.legend.Legend at 0x7fee10152748>



Neat! Now let's evaluate our model accuracy on the entire dataset. The predicted class label for a given input image can be computed by looking at the output of the neural network model and computing the index corresponding to the maximum activation. Something like

```
predicted_output = net(images) _, predicted_labels = torch.max(predicted_output,1)
```

```
predicted_output = model(images)
print(torch.max(predicted_output, 1))
fit = Loss(predicted_output, labels)
print(labels)
```

```
↳ torch.return_types.max(
  values=tensor([1.0211, 1.1970, 0.6002, 1.0526, 1.4884, 0.8140, 1.0334, 0.9941, 1.
    1.3968, 2.0169, 1.2345, 1.7619, 0.5816, 1.1849, 1.8226, 0.7461, 1.2836,
    1.6389, 1.8291, 1.8179, 1.2951, 1.2433, 1.5005, 0.9930, 0.4443, 1.2734,
    1.8999, 0.7364, 1.6322, 1.9442, 1.4469, 1.5511, 0.7837, 1.7236, 0.7170,
    0.9840, 0.8445, 1.0878, 1.7428, 1.3926, 1.1502, 1.7046, 0.9351, 1.1246,
    1.8468, 1.0007, 1.4352, 1.4131, 0.9919, 1.3044, 1.0443, 1.6108, 2.0653,
    1.1146, 1.8927, 1.3643, 1.6146, 1.7828, 0.6247, 0.9762, 1.4032, 0.6954,
    0.4679]), grad_fn=<MaxBackward0>),
  indices=tensor([7, 4, 4, 7, 0, 3, 7, 7, 4, 7, 9, 0, 1, 4, 4, 2, 3, 4, 9, 9, 4, 0
    9, 8, 3, 4, 4, 2, 1, 1, 1, 4, 4, 4, 8, 4, 7, 4, 3, 0, 7, 7, 7, 4, 4, 7,
    4, 4, 7, 2, 7, 9, 8, 4, 1, 4, 4, 4, 0, 4, 7, 4]))
  tensor([8, 4, 8, 5, 6, 3, 5, 7, 6, 7, 9, 0, 1, 2, 6, 2, 3, 6, 9, 9, 4, 0, 4, 4,
    5, 8, 0, 4, 6, 2, 1, 1, 1, 6, 2, 6, 2, 6, 7, 4, 3, 0, 7, 7, 7, 4, 6, 7,
    6, 4, 7, 2, 9, 9, 8, 4, 1, 2, 2, 8, 0, 4, 5, 3]))
```

```
def evaluate(dataloader):
    # Q4.6 Implement a function here that evaluates training and testing accuracy.
    # Here, accuracy is measured by probability of successful classification.
    correctPredictions = 0
    totalNumLabels = 0
    c = 0
    for images, labels in dataloader:

        images = images.view(-1, 28*28).requires_grad_()
```



```

        # Forward pass
        outputs = model(images)
        # Get predictions from the maximum value
        _, predicted = torch.max(outputs.data, 1)
        # Computing the total number of labels
        totalNumLabels += labels.size(0)
        # Computing the total number of correct predictions
        correctPredictions += (predicted == labels).sum()
    correctPredictions = correctPredictions.numpy()
    FinalAccuracy = correctPredictions*100/totalNumLabels
    return FinalAccuracy

print('Train acc = %0.2f, test acc = %0.2f' % (evaluate(trainDataLoader), evaluate(testDataLoader)))

☞ Train acc = 60.20, test acc = 59.27

```